HAPPIEST HEALTH · DIGITAL TWIN CHALLENGE 2026

GlucoTwin

A continuously updated virtual Type 2 diabetes patient that estimates, up to 2 hours ahead, the risk that a post-meal glucose excursion is about to begin, and shows a doctor why.

Team: team mono · **College / incubator:** Amrita Vishwa Vidyapeetham, Chennai
Members: Mohnish Shanmugavadivel (sole member)

Research and educational prototype on **synthetic patients only**. Not a medical device. Outputs are not diagnoses or treatment instructions.

Problem and healthcare use case

Condition: Type 2 diabetes (India)

- Post-meal glucose excursions are common in T2D, and a clinician reviewing CGM usually sees them **after** they happen.
- The challenge asks for a twin that fuses **simulated EHR** with **wearable / CGM streams** and predicts an adverse event, e.g. "a sudden glucose spike 2 hours in advance".

Use case

- A doctor monitors several remotely monitored patients.
- The twin flags who is likely to cross 180 mg/dL in the next 2 h, while glucose is still normal, with a forecast band and the reasons.

Event definition (frozen before modelling)

At time t, asked only if the latest CGM ≤ 180 mg/dL. **Positive** if ≥ 3 consecutive 5-min CGM readings > 180 mg/dL (≥ 15 min; gaps break the run) **begin** within (t, t + 120 min].

Why "onset": predicting "above 180" while already above 180 is trivial and useless clinically.

Outputs per CGM reading

- Glucose forecast at 30/60/90/120 min with an 80% interval
- Probability of excursion onset within 2 h, alert at a validation-chosen threshold
- Grouped explanation (associations, not causes)

Why a digital twin, and not a single risk score?

Continuously updated

A versioned per-patient state is updated with **every** observation: glucose and trend, 24 h time-in-range, meals and carbs on board, activity, data quality. Forecasts are re-issued at each reading.

Fuses record + streams

The static record (labs, medications, history) gives context and a prior. Streams give the current state. Both feed the same leakage-safe feature builder offline and online.

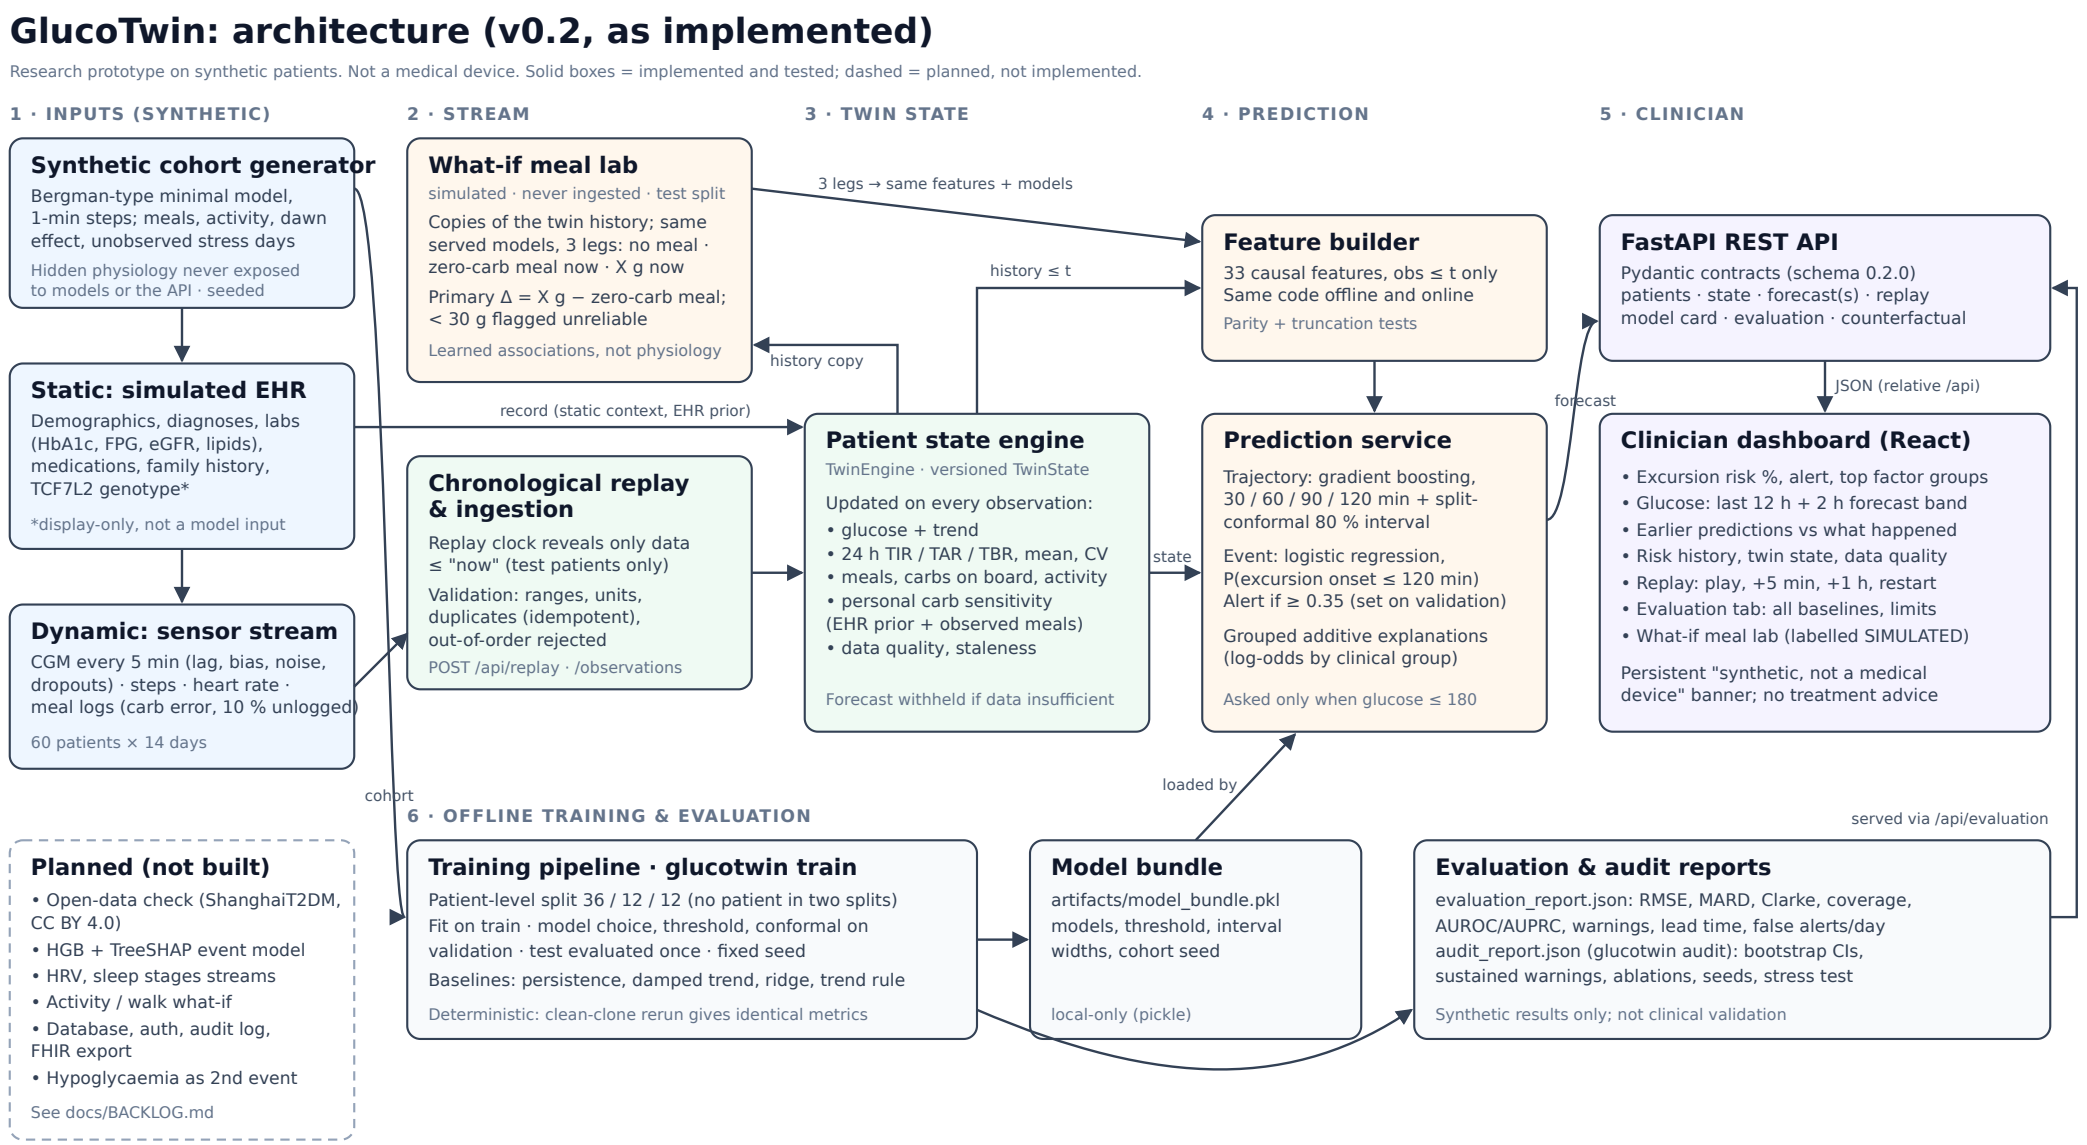
Personalised

Each twin learns its patient's **carb sensitivity** (mg/dL per g) from observed meal responses, shrunk towards a record-based prior. This changes the expected post-meal rise.

What it is not (yet)

The twin does not simulate physiological interventions such as activity ("what if she walks after lunch?"). That would need per-patient physiological parameter fitting, which is on the roadmap, not built. Its what-if meal lab only re-runs the learned forecasters on a hypothetical meal (an association, not a mechanism). Our twin is a **state-tracking forecaster** with one personalised physiological parameter.

Architecture (as implemented)



Data: synthetic, by rule and by necessity

- Rules: "strictly required to use anonymized, open-source, or synthetic datasets".
- Real CGM + EHR datasets with 5-min data (OhioT1DM, T1DEXI, DiaTrend) need data use agreements, which is not feasible before the deadline.
- **Generator:** a Bergman-type minimal model with meals, activity, dawn effect and unobserved "stress days". **Hidden** physiology is never shown to models.
- **Observable:** record (labs derived with noise), CGM with lag, bias, AR noise and dropouts, steps, heart rate, meal logs with carb error and 10% unlogged.
- 60 patients × 14 days; 5-min CGM; seed 20261017.

Plausibility check	Ours	ShanghaiT2DM*
Time in range 70–180	77.8%	77.7%
Time above 180	21.8%	20.0%
Time below 70	0.4%	2.4%

*Published summary of an open Chinese T2D cohort (Zhao et al., Sci Data 2023), quoted from a secondary copy. This is a plausibility check, not validation.

Known gaps. Lows are under-represented: hypoglycaemia only occurs on sulfonylurea or insulin in our simulator. Synthetic HbA1c is derived from the evaluated 14 days (corr 0.95 with that window's mean CGM). The ablation shows it does not inflate results (slide 8).

Evaluation protocol and leakage safeguards

Protocol

- **Patient-level split** 36 / 12 / 12; no patient appears in two splits.
- All models are fitted on train only. **Model choice, alert threshold and conformal intervals use validation only.** Test is evaluated once.
- Baselines on the identical test rows: persistence, damped trend, ridge (trajectory); trend rule (event).
- Deterministic: a **clean clone** with a fresh environment reproduced every metric exactly (wave-2 audit).

Leakage safeguards (tested)

- Features use observations ≤ t only. A **truncation test** shows that removing the future never changes past features.
- The **same code** builds features offline and online (parity test).
- Personal carb sensitivity uses a meal only after its 2-h response window has ended.
- **Mutation check:** injecting a 30-min future shift makes two leakage tests fail.
- The dashboard replay can only show held-out test patients and never reveals data after "now" (API test).

Results on 12 held-out synthetic patients

reproduced in a clean clone

Glucose forecast RMSE (mg/dL)

Model	30 min	60 min	120 min
Persistence (last value)	18.7	31.4	46.6
Damped trend	17.4	31.1	50.4
Ridge	11.2	17.0	27.9
Gradient boosting (served)	**10.4**	**17.2**	**25.9**

95% patient-bootstrap CI, served model: 30 min 9.8–11.0; 120 min 23.2–29.1. 80% conformal interval coverage on test: 0.76–0.80. Clarke A+B ≥ 95.7%.

Excursion onset within 120 min

Model	AUROC (95% CI)	Sens / Spec	False alerts / day
Trend rule baseline	0.655 (0.583–0.736)	0.44 / 0.74	6.3
Logistic regression (served)	**0.884 (0.839–0.923)**	**0.67 / 0.90**	**1.4**
Gradient boosting	0.936 (0.915–0.955)	0.78 / 0.90	1.1

Test prevalence 14.6% of eligible time points. Logistic regression is served because its explanation is exact and additive. Gradient boosting discriminates better (roadmap: serve it with SHAP).

These describe our simulator and are an **upper bound**. Real-data studies report a 30-min RMSE of about 17–20 mg/dL. Not clinical validation.

Source: reports/evaluation_report.json, reports/audit_report.json

A stricter look: our own audit

wave-2 · glucotwin audit

How often does the doctor get warned? (served model)

Definition	Share of onsets
Any alert in the 2 h before onset	91% (0.83–0.98)
…with ≥ 30 min lead	87%
Sustained: alert still on just before onset	**69% (0.48–0.83)**

Median lead: any alert 95 min; sustained 85 min (capped at 120). False-alert episodes per monitored day: 1.36 (0.92–1.92). The headline "91% warned" uses the lenient definition, so we report both.

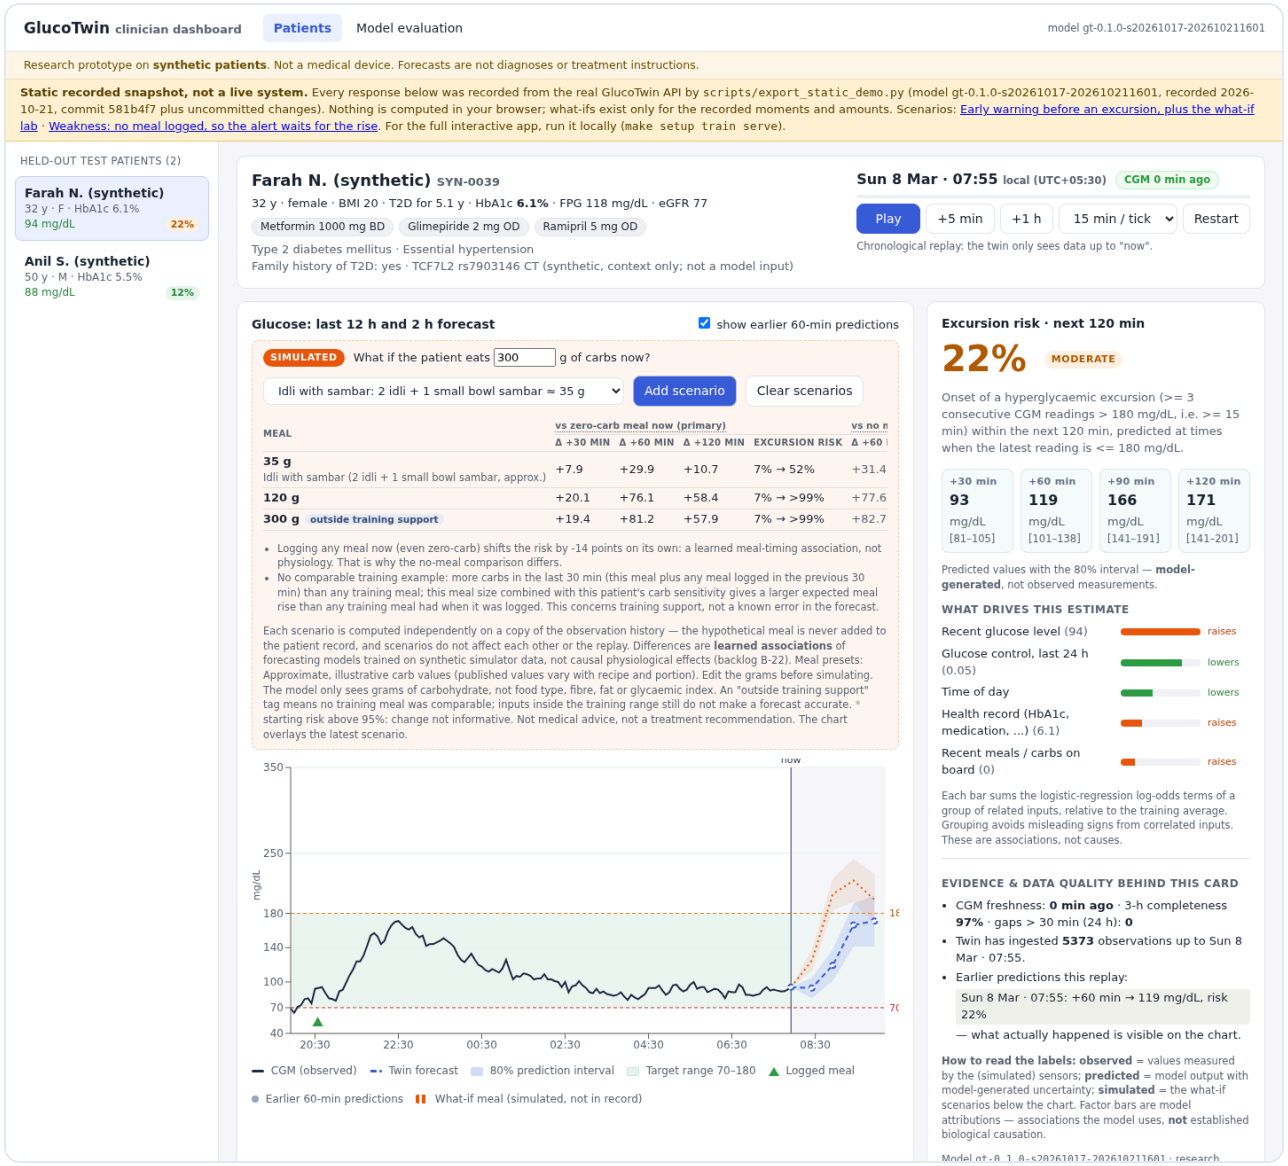
Robustness

Check	Logistic AUROC	Sustained warned
Main cohort	0.884	69%
3 other cohort seeds	0.883–0.910	72%–78%
No record (EHR) features	0.907	73%
No meal features	0.849	68%
Stress: 60% logged, ±50% carbs, 0–45 min late, noisier CGM	0.904	70%

Stress test: RMSE at 30 min rises from 10.4 to 15.4; warnings with ≥ 30 min lead fall from 87% to 78%. Direct record features did **not** help held-out discrimination; the record contributes through the personal prior and context.

Source: reports/audit_report.json (deterministic; 1000 patient-level bootstrap resamples)

Clinician dashboard: what the doctor sees



- **Patient list:** held-out test patients, latest glucose, latest risk.
- **Record header:** labs, medications, history.
- **Glucose chart:** last 12 h, 2-h forecast with 80% band, meals, earlier predictions versus what happened.
- **Risk card:** probability (never shown as 100%), alert wording "Consider reviewing", top factor groups.
- **Risk history** and **twin state:** TIR, carbs on board, learned carb sensitivity, data quality.
- **Replay:** the twin only sees data up to "now".
- **Evaluation tab:** every baseline and limitation.
- **What-if meal lab** (simulated): *X* g eaten now versus a zero-carb meal at the same time, plus a grey no-meal comparison. Learned model associations, not physiology; meals under 30 g flagged unreliable.

Two demo moments: an early warning and a late one

SYN-0039 · Sun 8 Mar (held-out patient)

- 08:53 a 140 g breakfast is logged; at 08:55 glucose is **89 mg/dL**.
- Risk is **>99%**, the alert switches on, and the forecast reaches about 205 mg/dL at +60 min.
- Onset at 10:20, peak about 207 mg/dL. The alert stayed on throughout: an **85-min sustained warning**.
- Honest detail: the rise came later than the 30-min forecast (103 observed versus 128 forecast at 09:25).

Deep link /?patient=SYN-0039&start=151.833 · verified against the live API

SYN-0019 · Thu 12 Mar (held-out patient)

- 14:45: glucose 88 mg/dL, risk 31%, just **below** the 35% threshold. **No meal logged.**
- 14:50: glucose rising fast, risk 36%: the alert fires, only 45 min before the first reading above 180 (15:35); peak about 191. **Warned, but late.**
- Unlogged meals are the main failure mode: removing meal features costs the most AUROC, and the stress test cuts lead time.

Deep link /?patient=SYN-0019&start=253.667 · Of 298 test onsets, 205 (69%) had a sustained warning and 272 had an alert at some point in the prior 2 h. scripts/find_demo_scenario.py lists 53 daytime onsets (08:00–21:59, ≥ 30 h into the data) whose alert was off at the last reading before onset.

Safety boundaries and limitations

Safety boundaries (built in)

- A "synthetic, not a medical device" banner on every screen; synthetic IDs (SYN-####) and names tagged " (synthetic)".
- No instructions: the alert text is "Consider reviewing", never "give insulin" or "eat".
- Explanations and what-if meal results are labelled as associations, not causes; a what-if is never written to the record.
- No risk number when data are insufficient or stale; no onset risk when glucose is already above 180.
- Probabilities are never displayed as 100% or 0%.
- Local demo: no authentication, so it is **not for real patient data**.

Limitations

- Evaluated only on our own simulator; no real-patient or external validation.
- 12 test patients, so the confidence intervals are wide.
- Meal logging is near-real-time in the main cohort, which is optimistic.
- Lows are under-represented (0.37% of readings < 70 mg/dL vs 2.4% published); the low-glucose watch is a lower bound, **not a probability**.
- Vs simulator ground truth, the What-If underestimates late meal effects (+120 min MAE 29.0 vs 12.8 mg/dL for a cohort-average curve).
- Personalisation is shown only in simulation (hidden meal gain recovered, r 0.95 by day 3; +30 min RMSE −0.52 mg/dL, 95% CI 0.12–0.96).
- No HRV or sleep streams; genotype is synthetic and display-only. No regulatory claims (CDSCO SaMD, DPDP, HIPAA).

Potential impact and the path to real validation

If validated, the intended value is

- Earlier, reviewable warnings for remotely monitored T2D patients, in place of after-the-fact CGM review.
- A per-patient view that explains *why* risk is high (meal, trend, recent control).
- A transparent evaluation template (baselines, CIs, alert burden) for any future model.

We have run no pilot and have no user or outcome data. No impact figures are claimed.

Validation roadmap (planned)

1. Retrospective check on open data (ShanghaiT2DM, CC BY 4.0; 15-min CGM).
2. Serve gradient boosting with SHAP; within-patient chronological evaluation.
3. Clinician usability review of the dashboard (think-aloud, no patients).
4. Ethics-approved retrospective study on consented, de-identified data.
5. Prospective **silent-mode** study (alerts not shown to clinicians), then a regulatory pathway assessment.

Engineering and reproducibility

Stack

- Python 3.11 · NumPy · pandas · scikit-learn · Pydantic v2 contracts · FastAPI
- React 19 + TypeScript + Recharts (Vite)
- pytest (151 tests) · ruff · GitHub Actions CI · pinned requirements.lock / package-lock.json

Reproduce

```
make setup
make train     # same metrics every run
make serve     # http://localhost:8000
glucotwin audit  # CIs, ablations, stress test
make check     # lint + tests + build
```

Verified in the wave-2 audit

- Clean clone + fresh environment: tests pass and the report is identical.
- Audit re-implementation matches every headline metric exactly.
- pip-audit and npm audit: no known vulnerabilities (after a pytest upgrade).
- Fixed: API returned 500 on extreme replay offsets (now 422); training patients could be replayed (now refused).
- Open source: MIT licence.

Thank you

GlucoTwin: a transparent, reproducible digital-twin prototype for early warning of post-meal hyperglycaemia.

Team: team mono · Amrita Vishwa Vidyapeetham, Chennai · Mohnish Shanmugavadivel

All results are from synthetic patients. Not a medical device. Not clinical validation.